

DINC CETINTAS

AI Engineer

EU Citizen - no visa sponsorship required | Cyprus | dincetintas24@gmail.com | +974 51444964 |
linkedin.com/in/dinc-cetintas-79a04b205 | GitHub

PROFILE

AI Engineer who designs, builds, deploys, evaluates, and scales production AI systems in regulated financial environments. I own LLM and RAG platforms end-to-end - architecture, evaluation, agentic orchestration, APIs, model serving, and observability - delivering reliable, auditable, low-latency products used across 7+ business teams and 200+ internal users. I combine software engineering, MLOps, and GenAI to turn mission-critical banking workflows into measurable gains in cost, speed, and risk reduction.

PROFESSIONAL EXPERIENCE

AI Engineer - Commercial Bank of Qatar, Doha

Jan 2024 - Present

- Architected and owned production LLM and RAG platforms for document intelligence, workflow automation, and conversational AI, deployed across 7+ business teams and 200+ internal users at a regulated bank.
- Designed the RAG architecture - hybrid retrieval, semantic chunking, reranking, and citation grounding over 900+ compliance documents - delivering auditable, source-linked answers at 87% retrieval accuracy.
- Built automated evaluation pipelines measuring retrieval quality, grounding, and hallucination rate, with human review loops, gating every release on quality thresholds before production.
- Productionized agentic workflows (LLM reasoning, tool calling, OCR, structured outputs) powering inquiry automation - 500+ requests/week, 70% manual-load reduction, COO Problem Solver Award - and credit-risk & financial-analysis report generation.
- Built and owned the FastAPI / REST services and model-serving layer (Docker, Kubernetes, MLflow) with production monitoring and observability, cutting inference latency ~40% and ensuring reliability.
- Designed and deployed a production fraud-detection pipeline combining engineered risk features, model evaluation, and monitoring, achieving ROC-AUC 0.93 while strengthening real-time transaction risk controls.
- Delivered an on-prem Arabic<->English document-translation platform that preserves layout and keeps sensitive data fully in-house, eliminating ~\$120K/yr in third-party translation spend.

Software Engineer - MeeApps, Istanbul

Feb 2023 - Jan 2024

- Architected and delivered a high-frequency trading platform for 23 brokerage firms managing \$10M+ in portfolios, achieving microsecond-level execution by optimizing ITCH, OUCH, and TCP protocols for ultra-low latency.
- Built and owned the full production stack - React/C#/C++ frontend, backend services, and co-located live market-data feeds - under strict real-time reliability constraints.
- Owned end-to-end system performance and reliability, sustaining 99.9% uptime under live trading load.

Machine Learning Engineer - Schneider Electric, Istanbul

Jul 2022 - Dec 2022

- Built and deployed demand-forecasting and stock-out-risk models that reduced inventory gaps 12-15% during COVID supply shocks, directly improving fulfillment reliability.
- Engineered features from sales, supply-chain, and distributor data and built reproducible model pipelines with rigorous evaluation (cross-validation, holdout testing) for scalable deployment.
- Partnered with logistics stakeholders to define success metrics and operationalize forecasts into fulfillment planning.

EDUCATION

BASc, Electronics Engineering - Sabanci University, Istanbul - 2018 - 2023

Exchange Semester, Electrical & Electronics Engineering - Universitat Politecnica de Valencia (UPV), Spain - 2022

TECHNICAL SKILLS

AI Systems / GenAI: LLM systems, RAG (hybrid retrieval, reranking, citation grounding), agentic & multi-step workflows, tool / function calling, prompt & context engineering, AI orchestration, structured outputs

AI Engineering / LLMops: Automated evaluation (retrieval quality, grounding, hallucination mitigation), human review loops, benchmarking, observability & logging, latency optimization, reliability engineering

Production / Platform: FastAPI / REST APIs, backend & service architecture, model serving, Docker, Kubernetes, MLflow, monitoring

ML Foundations: Supervised & unsupervised learning, feature engineering, time-series forecasting, A/B testing, statistical evaluation

Cloud & Data: Azure (OCR, Translation, Custom Models), AWS, SQL

Programming: Python; full-stack (React, C#, C++)

SELECTED AI SYSTEMS

- **Enterprise RAG Platform** - 900+ compliance documents; auditable, source-linked answers at 87% accuracy
- **AI Email / Inquiry Assistant** - 500+ requests/week automated; 70% manual-load reduction
- **Credit Risk Report Generation Platform** - agentic generation of analyst-ready credit-risk reports
- **Financial Analysis Agent Workflows** - multi-step LLM reasoning + tool calling over financial data
- **OCR & Document Intelligence Platform** - 7 extraction pipelines for Arabic/English ID & financial docs
- **Arabic<->English Translation Platform** - on-prem, layout-preserving; ~\$120K/yr third-party spend eliminated

CERTIFICATIONS & COURSES

Operationalizing LLMs on Azure (Duke) | Generative AI Applications with RAG & LangChain (IBM) | AI Agents (IBM) | Generative AI Techniques & Applications (Duke) | Machine Learning Specialization (Stanford) | TensorFlow in AI/ML/DL (DeepLearning.AI) | Deep Learning for Sequences & Time Series (DeepLearning.AI) | Visa Data Science Bootcamp

ADDITIONAL

Languages: English | Links: LinkedIn, GitHub