DINC CETINTAS

SENIOR LLM DATA SCIENTIST

EU Citizen, Cyprus | dinccetintas24@gmail.com | +974 51444964 | LinkedIn | GitHub

PROFILE

Senior LLM Data Scientist with hands-on experience in LLM fine-tuning, evaluation, and data-pipeline engineering in regulated financial environments. I turn raw, unstructured enterprise data into training- and retrieval-ready datasets (cleaning, filtering, deduplication), fine-tune and evaluate transformer models (GPT, T5, BERT), and build RAG systems over vector databases - all in Python. I partner closely with product and data teams to translate requirements into working technical designs, take strong ownership of outcomes, and explain complex trade-offs clearly to executives and engineers alike.

PROFESSIONAL EXPERIENCE

LLM Data Scientist / AI Engineer - Commercial Bank of Qatar, Doha *Jan 2024 - Present*

- Built data pipelines that turn raw, unstructured enterprise data into training- and retrieval-ready datasets - cleaning, filtering, deduplication, and structured extraction across 900+ compliance and financial documents.
- Fine-tuned and evaluated transformer models (GPT, T5, BERT), refining datasets and prompts against automated quality benchmarks (retrieval quality, grounding, hallucination) with human review loops before every release.
- Designed RAG systems over vector databases - hybrid BM25 + vector retrieval, reranking, citation grounding, and context engineering (tokenization-aware chunking and context-window budgeting) - reaching 87% source-linked accuracy over 900+ documents.
- Built a reusable AI platform (shared retrieval/embeddings service, evaluation framework, and model-serving API) consumed by 7+ teams, cutting time-to-delivery for each new use case.
- Partnered with product managers, data specialists, and business stakeholders to understand requirements and roadmaps, translate business concepts into technical designs, and feed insights back into the roadmap - owning use cases end-to-end across multiple squads.
- Shipped clean, maintainable Python integrating models into live banking systems (FastAPI/REST serving), with production monitoring and feedback loops that detect regressions and route them back into evaluation and dataset refinement.
- Developed production LLM workflows for credit-risk and financial analysis - retrieval, tool calling, LLM reasoning, and JSON-schema-validated outputs - producing structured, analyst-ready reports.
- Designed and deployed a fraud-detection model (ROC-AUC 0.93) with engineered features, evaluation, and monitoring, strengthening transaction-risk controls.

Software Engineer - MeeApps, Istanbul *Feb 2023 - Jan 2024*

- Built a low-latency order-execution platform - smart order-routing and execution algorithms - adopted by 23 brokerage firms managing $10M+ in portfolios, achieving microsecond-level execution.
- Delivered the full production stack (React/C#/C++ frontend, backend services, and live market-data feeds) in a cross-functional team, optimizing ITCH, OUCH, and TCP protocols against co-located exchange servers for ultra-low latency.
- Owned end-to-end system performance and reliability, sustaining 99.9% uptime under live trading load.

Data Scientist - Schneider Electric, Istanbul *Jul 2022 - Dec 2022*

- Built and deployed demand-forecasting and stock-out-risk models in Python (pandas, scikit-learn) that reduced inventory gaps 12-15% during COVID supply shocks.
- Engineered features from sales, supply-chain, and distributor data and built reproducible pipelines with rigorous evaluation (cross-validation, holdout testing) for scalable deployment.
- Partnered with logistics stakeholders to define success metrics and operationalize forecasts into fulfillment planning.

EDUCATION

BASc, Electronics Engineering - Sabanci University, Istanbul - 2018 - 2023

Exchange Semester, Electrical & Electronics Engineering - Universitat Politecnica de Valencia (UPV), Spain - 2022

TECHNICAL SKILLS

LLM & Model Development: Fine-tuning (SFT), evaluation & dataset refinement, transformers (GPT / T5 / BERT), tokenization, context engineering, prompt versioning

Data / ML Engineering: Training- & retrieval-ready data pipelines (cleaning, filtering, deduplication, structured extraction), feature engineering, RAG, vector databases

Evaluation / LLMOps: Automated evaluation (retrieval quality, grounding, hallucination), human review loops, production monitoring & feedback loops, model routing & fallback strategies

Languages & Tools: Python (pandas, scikit-learn), SQL, FastAPI, Docker, Kubernetes, MLflow; Azure, AWS

Domain: Financial services, regulatory compliance (auditability, traceability), risk

SELECTED AI SYSTEMS

- Enterprise RAG Platform - hybrid retrieval + reranking over 900+ docs on a vector DB; 87% source-linked accuracy

- Enterprise Data Pipelines (Training & Retrieval) - cleaning, filtering, deduplication & structured extraction into model-ready datasets

- Model Evaluation Framework - retrieval-quality, grounding & hallucination benchmarks; human review; release gating

- Credit Risk & Financial-Analysis LLM Workflows - tool calling + LLM reasoning + schema-validated, analyst-ready reports

- Real-Time Fraud-Detection Model - ROC-AUC 0.93; engineered features, evaluation, monitoring

CERTIFICATIONS & COURSES

Operationalizing LLMs on Azure (Duke) | Generative AI Applications with RAG & LangChain (IBM) | AI Agents (IBM) | Generative AI Techniques & Applications (Duke) | Machine Learning Specialization (Stanford) | TensorFlow in AI/ML/DL (DeepLearning.AI) | Deep Learning for Sequences & Time Series (DeepLearning.AI) | Visa Data Science Bootcamp

ADDITIONAL

Languages: English (professional)